

How the Coach reviews

What the Coach reads, which sciences it brings, how it decides what to say, how it scores, and how it suggests and fixes copy. Written for the market teams who use it and the people who extend it.

Acquisition Journey Builder · September 2026


One principle above everything

Evidence beats opinion. The sciences say what may influence people in general. Only DAZN's own numbers say what happens in a DAZN journey. When the Coach has neither, it says so. "I don't have enough evidence to recommend changing this" is a strength, not a failure.

1. What the Coach is

The Coach is an expert reviewer of a whole acquisition journey: landing page, plan choice, how to pay, sign-in, account, ZIP check, checkout, confirmation. It does not redesign screens and it does not invent optimisation ideas. It looks at what is on screen, says what it sees, says what a science suggests that may do, and recommends a change only when something supports it.


Every finding follows one chain:



Example. Observation: "Pay now" (\$279.99/year) is preselected. Evidence: Default Effect, a peer-reviewed science. Interpretation: the preselected option gets chosen far more often; there is no recorded reason for giving it that advantage. Recommendation: none, not enough evidence. Confidence: Likely. The Coach then says how to become sure: look at how the two ways to pay split today, and test before changing the default.

2. The layers, from goal to score

The review runs through fixed layers. Each one only receives what the one before produced, so a brain cannot see anything that is not on screen, and the score cannot count anything the Coach did not first check for evidence.



2.1 The brief

Before it looks at a single screen the Coach is told what the journey is for. Market and journey come from the panel. The goal and, when it applies, the package to prioritise, come from the goal sheet. The ten baseline questions run whatever the goal is; the goal only changes their weight and adds the checks that belong to it.

2.2 The facts

The journey is read through the same rules the product renders with. A price is the formatted price on the card, a badge is the derived ribbon, the default is the option whose radio is filled, the copy is every string on every screen. Nothing is judged at this layer. Because the brains only get facts, an observation is always something you can see by opening the screen.

2.3 The brains

Seven specialists, each owning a part of the knowledge (section 4). A brain writes findings in the chain above and nothing else. It cannot change a score or override another brain.

2.4 The Coach

The Coach sits above the brains and does three things a committee cannot. It **merges** two brains reading the same element into one finding with both sciences. When they **disagree** about whether to recommend anything, it keeps the observation, withholds the recommendation and says so. And it applies the **evidence guard**: a fix must rest on the content itself or on DAZN evidence. Science alone can raise a question, never demand a change, and nothing science-only is ever marked Certain.

3. The Acquisition Baseline: ten questions, always asked

These run on every review, whatever the goal, like truth rules. A goal can make one matter more (marked with a small triangle in the panel); it never switches one off. Each finding answers exactly one of them, and each question is scored on its own.

#	Question	Examples of what pulls it down
1	Completion. Can the user successfully reach subscription?	A screen with no primary action; a step that blocks the way to payment.
2	Decision clarity. Can they understand what they are choosing?	Placeholder text left on screen; slips like “1 locations”; a price with no reference beside it.
3	Package comprehension & comparison. Are differences clear?	Identical feature lines on every card; a plan described only by naming other plans.
4	Decision friction. Is unnecessary effort getting in the way?	Marketing consent inside account creation; the ZIP check after the account is made.
5	Purchase confidence. Do they understand what they get and what happens next?	A bill that says /month for a yearly amount; two renewal dates; one phrase meaning two things at checkout.
6	Offer comprehension. Are offers and trials communicated correctly?	A yearly saving that exists in the numbers but is never said; a goal about a trial when no trial is shown.
7	Journey consistency. Does each step keep the promise made earlier?	Landing names teams the headlined plan does not include; later screens written for one plan only; “sign in” here, “log in” there.
8	Informed choice & trust. No misleading defaults, hidden conditions, unsupported claims, artificial pressure?	“Most Popular” with nothing behind it; a preselected option with no recorded reason; “free” on the way to a paid plan.
9	Commercial alignment. Does the journey support the stated business objective?	The screen recommends a plan other than the one you asked to prioritise; the goal is annual but monthly is the default.
10	Evidence integrity. Are the Coach's own conclusions supported?	Drops for every reading the guard had to rein in. A judgement of the Coach, not of the journey.

4. Goals: what the Coach is allowed to optimise

A goal names a lever the business may pull. It raises the weight of the baseline questions it depends on and adds checks of its own. Goals that need a target (which content, which offer) ask for it in the sheet; the Coach then traces that word through every screen.

Goal	What the Coach checks	Raises
Drive Specific Package	Is the named package first, recommended, and the one the copy points at? A recommendation on another plan is a fix.	Clarity, comparison, alignment, trust
Drive Annual Plan	Is annual offered, preselected, and does the screen say why a year is worth committing to? A monthly default with this goal on is a fix.	Purchase confidence, offers, alignment, trust
Acquire for Specific Content	Where the named content appears, screen by screen, and where it goes cold.	Journey consistency
Acquire Specific Audience	Whether any screen speaks to that audience. Fit is a question for research, and the Coach says so.	Decision clarity
Drive Specific Offer	Whether the named offer is on screen at all, and carried through. Never its economics.	Completion, confidence, offers, alignment, trust
Drive Specific Benefit	Whether the benefit is on the cards, whether it separates plans or is shared by all, and whether the landing page leads with it.	Clarity, comparison
Drive Bundle / Add-on	Whether a bundle exists in the set and whether position and emphasis favour it.	Comparison, offers, alignment, trust
Drive Immediate Start	How many screens precede payment, whether anything says when watching starts, and any invented urgency words (a fix).	Completion, friction, confidence
Content-to-Package Match	Which plan carries every team the landing page names, confirmed by the plan's own team list, and whether the screen favours it.	Comparison, consistency, alignment
Campaign-to-Journey Continuity	Where the campaign proposition, in its own words, is and is not visible from entry to checkout.	Journey consistency

5. The seven brains

Knowledge is split into specialists so that no single prompt or rule set can justify anything with some theory. Each brain reads the same facts and writes findings only in its own domain.

Brain	What it reads	Typical finding
Decision	How people compare options on several attributes at once.	Every plan lists the same four features, so only teams and price differ; the lines cannot help the comparison.
Choice architecture	Order, defaults, the recommended plan, how many options, grouping.	"Pay now" is preselected; the priciest plan sits in the middle and is the recommended one.
Cognitive & clarity	Mental load per screen, one name for one thing, information arriving when needed.	Five screens share the title "Choose your subscription"; "Pay now" means a cadence on one screen and paying on the next.
Trust & risk	Claims, contradictions and conditions near payment.	Checkout shows \$279.99 per month; Cadence shows the same amount per year. Two renewal dates a few lines apart.
Journey / UX	The journey as one thing: promise, expectation, transition, friction with or without purpose.	The landing page promises Yankees and Nets under the MSG+ headline; MSG+ does not include them.
Commercial	The goal you set, checked against the journey. Targets traced screen by screen.	"Knicks" is on Landing, Subscription and Confirmation, but not on Cadence, Login or Checkout.
Copy	Does the copy say what the reader gets, is one action one name, is it finished. Plus the AI reading when a key is set.	"Sign in" on Landing, "Log in" on Login and Zip code. A lower-case word after a full stop on the confirmation.

6. The science stack

A catalogue of fifty theories was considered. Twenty were kept as the core: enough to understand a decision journey, few enough that a reading cannot be justified by whatever theory happens to fit. Three methods sit above them all. A handful of named extensions are allowed where the brief calls them out. Each science has a strength: a strong one earns Likely on its own; a contested one earns only Possible. No science earns Certain.

Decision and choice

Science	What it governs	Brain
Judgment and Decision Making	Overall human decision-making and evaluation of alternatives	Decision
Choice Architecture	How presentation and organisation of choices influence decisions	Choice
Default Effect	Effect of preselected options	Choice
Choice Overload (<i>contested</i>)	Effects of having too many alternatives	Choice
Multi-Attribute Decision Making	How people evaluate options differing across several attributes	Decision

Science	What it governs	Brain
Prospect Theory	Gains, losses, risk and reference-dependent decisions	Decision
Framing Effect	How equivalent information changes decisions depending on presentation	Decision
Anchoring Effect	How an initial value affects later judgments	Decision
Compromise Effect (<i>contested</i>)	Tendency to prefer an intermediate option	Choice
Mental Accounting	How consumers mentally categorise costs and value	Decision
Reference Price Research	How a price is judged against an expected or comparison price	Decision

Understanding and journey

Science	What it governs	Brain
Cognitive Load Theory	Mental effort required to understand a screen or decision	Clarity
Processing Fluency	Ease with which information is processed	Clarity
Hick-Hyman Law	Number and complexity of choices against decision time	Clarity
Information Foraging Theory	How users follow cues toward what they want	Journey
Mental Models	Whether the journey matches expectations of how such things work	Journey
Expectation-Confirmation Theory	The gap between what users expect and what the experience delivers	Journey

Trust and persuasion

Science	What it governs	Brain
Persuasion Knowledge Model	How people recognise and respond to persuasion attempts	Trust
Psychological Reactance Theory	Resistance when users perceive pressure or manipulation	Trust
Perceived Risk Theory	Uncertainty about financial, product or performance consequences	Trust

Named extensions and methods

Science	What it governs	Kind
Signaling Theory	How labels, prominence and recommendations communicate	Extension
Social Proof	Why claims such as "Most Popular" influence choice	Extension
Progressive Disclosure	Revealing information when it is needed	Extension
Goal-Gradient Hypothesis (<i>contested</i>)	Motivation rising as someone approaches a goal	Extension
Behavioral Sludge	Unnecessary administrative or behavioural friction	Extension
Recognition Rather Than Recall	Keeping needed information visible instead of remembered	Extension

Science	What it governs	Kind
Experimental Design / A/B Testing	Whether a proposed change actually causes improvement	Method, above all
Causal Inference	Separating correlation from cause	Method, above all
Internal consistency	Two strings in the same journey cannot both be true	Method

7. Evidence and confidence

Every finding names what backs it, on a fixed ladder. Higher rungs outrank lower ones whenever two readings disagree.

Level	Evidence	What it can justify
0	The content itself	A contradiction, a placeholder, two names for one thing. Can justify a Fix and Certain .
1	DAZN experiment	A controlled test. Can justify a Fix when the team supplies it.
2	DAZN analytics	What happens, not why. Can justify a Fix; a test settles cause.
3	Market or user research	Explains; does not measure.
4	Peer-reviewed science	A Check , at most Likely .
5	HCI or industry principle	A Check, at most Likely.
6	AI inference	A Check or a Note, at most Likely. Never dressed as science.

Confidence in three words

Certain	The content itself shows it. Only level 0 evidence, or the team's own data, earns this.
Likely	Established science says so; DAZN data has not confirmed it here.
Possible	A contested science, or a reading two brains disagree on. Worth checking, not a conclusion.

The evidence guard

After the brains have spoken, the Coach checks every finding against the ladder. A Fix with nothing above level 3 behind it is held at Check and its confidence capped at Likely, with a note saying why. An AI finding can never be Certain. Each time the guard has to step in, the Evidence integrity score drops a little, because a brain overstated a reading before the Coach caught it.

What a toggle for evidence would change

Telling the Coach that experiments or analytics exist does not raise confidence: the Coach still has not read them. What changes is the **next step** on each finding: “you have analytics, look at how the two cadence options split before changing the default” instead of “no DAZN evidence yet, only a controlled test can settle this”.

8. From finding to fix

Fix	The content contradicts itself or is unfinished. Change it. Always backed by level 0 evidence.
Check	A science says this may matter. Evidence is needed before acting; the finding says how to get it.
Note	Worth knowing, including what is working. Costs no points.

Click to see it

Clicking a finding opens its screen and lights up the exact words it is about, on the phone preview and in the field that edits them. The Coach reads the page for those words rather than being wired into every

screen, so any component that shows the text takes part.

One-tap fixes

Where the content itself proves what is wrong, the finding carries a fix: a text change scoped to the one field it names. Current fixes include the wrong unit at checkout, two renewal dates, “Log in” to “Sign in” everywhere, “1 locations”, a lower-case word after a full stop, the casing of a ribbon, renaming the “Pay now” option “Pay yearly” to match “Pay monthly”, and “Continue” on the sign-in screen where nothing is confirmed yet. Fixes go through the same store methods the panel uses, so a market override stays an override. After a fix the rules run again on their own; the finding leaves the list and the score moves.

What is never auto-fixed

Prices. Defaults. Which plan is recommended. The order of plans or steps. Any claim. These get a recommendation in words and a “Show on screen”, and the decision stays with a person.

9. How the score is made

Each of the ten baseline questions starts at 100. Every finding that answers it takes a share of what is left: a Fix costs 22 percent of the remainder, a Check 9 percent, a Note nothing, and each is scaled by confidence (Certain 1.0, Likely 0.7, Possible 0.4). Because penalties compound, four small slips read as “several small things” rather than as zero.

The overall score is the weighted average of the ten. Every question weighs 1; a question the goal depends on weighs 2 or 3. A goal never removes a question. The per-goal score is the same ten numbers seen through that goal's weights, so it is never a different arithmetic. Per-screen scores use only the findings on that screen.

Score	Word	Meaning
85 to 100	Ready to ship	Nothing the content proves is wrong; remaining checks are hypotheses.
70 to 84	Nearly there	A few fixes; the structure holds.
50 to 69	Needs work	Several contradictions or a question at risk.
Under 50	Not ready	Readers will meet things that are wrong.

The panel shows the points each finding is worth to its question (a green +N), where each question would land if its fixes were done and if its checks were settled too, and how far the whole score has moved since the review began.

10. The AI reading

The rules find everything the content proves. The AI reading covers what needs judgement: does the copy say what the reader gets, does each step keep the last one's promise, is anything pushing rather than informing. It runs after the rules, only when a key is set (or, on a developer's machine, through the signed-in Claude command), and its results merge through the same Coach and the same guard.

What it is told

The brief: the goal, the package, the constraints. The facts: every screen's strings exactly as on screen. The rules' findings already made, so it does not repeat them. And the discipline: quote exactly, never predict a number, recommendation null when evidence is missing, confidence at most Likely, severity at most Check, at most eight findings, fewer well-grounded ones beat many.

What it may never do

Recommend raising, lowering or inventing a price. Recommend moving a default or a “most popular” label to another option. Treat a trade-off between two stated goals as a recommendation. Present inference as science. Everything it returns is level 6 evidence and shows an AI mark on the card.

How it is checked on arrival

Each returned finding is checked against the vocabulary: a known brain, a baseline question, sciences from the stack, a screen that exists in this journey. Anything else is dropped. Confidence and severity outside the allowed range are lowered, and the guard runs again after the merge.

11. How copy is suggested

The Coach suggests copy in three ways, from most to least certain.

Way	How it works	Example
Derived from the content	The right words already exist somewhere in the journey; the fix reuses them. Nothing is invented.	The yearly option is renamed “Pay yearly” because the other option is “Pay monthly”. The checkout unit becomes “year” because the cadence screen says so.
Recommended in words	When the change needs a decision, the Coach says what the copy should do and leaves the writing to a person.	“One line saying the account is the last step before MSG+.” “Say the saving next to the yearly price. The price itself does not change.”
Written by the AI, applied by hand	The AI reading may propose wording. It arrives as a Check with an AI mark; applying it is a separate click, through the same store methods a typed edit uses.	A plan description that names what the plan includes in its own words instead of naming other plans.

Rules copy always follows

One action, one name, everywhere. A label says what the screen does. A claim is either backed or removed. Urgency is only ever a true date or time. Prices are commercial facts, never rewritten. Copy is written in the market's own language and matches the voice already on the screens. And no em dashes anywhere in the tool.

12. What the Coach will not do

Invent or change a price	It judges how a price is shown, compared and positioned. Never whether it is the right price.
Move a default on its own	It names the advantage a default gives and asks what supports it. Changing it is a decision for a person, ideally after a test.
Redesign a screen	It works within the journey the other team built: content, order, defaults, emphasis, copy.
Add urgency	Immediacy must be true. Invented urgency is flagged as a fix, not suggested.
Optimise one metric blindly	Two stated goals pulling apart is a trade-off; the Coach says so and withholds the recommendation.
Pretend to be sure	When evidence is missing it says so. That is the point of the whole design.

13. Words in the panel

Fix / Check / Note	Severity, section 8.
Certain / Likely / Possible	Confidence, section 7.
+N (green)	Points the question gets back if this finding is resolved.
Small triangle beside a question	Your goal makes this question weigh more.
AI mark on a card	The finding came from the AI reading. Level 6 evidence.
Red or amber count beside a screen name	How many findings touch that screen; red if any is a fix.
Show on screen	Opens the screen and lights up the words the finding names.
Details	Where the Coach saw it, how to be sure, how sure it is, which brain and sciences read it.

Source documents: coach/docs/ACQUISITION_COACH_BRIEF.md, ACQUISITION_SCIENCE_STACK.md, ACQUISITION_BASELINE.md. Code: src/demo/coach/review.